


# Class 07 Lab

Darby Patterson

## Example of K-means clustering

First, we make up some data with known structure, so we know the suspected answer.

```
tmp <- c(rnorm(30, mean=-3), rnorm(30, mean= 3))  
x <- cbind(x= tmp, y= rev(tmp))  
plot(x)
```



Now we have structured data in x. Lets see if k-means is able to identify the groups.

```
k <- kmeans(x, centers=2, nstart= 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.092537	2.636453
2	2.636453	-3.092537

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 55.36825 55.36825
(between_SS / total_SS = 89.9 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Now we will explore k.

```
k$size
```

[1] 30 30


k\$centers

	x	y
1	-3.092537	2.636453
2	2.636453	-3.092537

```
k$cluster
```


[illegible]

```
plot(x, col= k$cluster)
```



Now we add the cluster centers:

```
plot(x, col= k$cluster)  
points (k$centers, col= 'purple', pch=15)
```



## Hierarchical Clustering

Lets use the same data stored in `x` and using the `hclust()` function.

```
clustering <- hclust(dist(x))  
clustering
```


Call:

```
hclust(d = dist(x))
```

```
Cluster method   : complete  
Distance         : euclidean  
Number of objects: 60
```

```
plot(clustering)  
abline(h= 10, col= 'pink')
```

## Cluster Dendrogram




To get our results ( membership vector), we need to “cut” the tree. using `cutree()`.

```
subgroups <- cutree(clustering, h=10)
subgroups
```

[illegible]

Plotting this...

```
plot(x, col=subgroups)
```



You can also “cut” your tree with number of desired clusters:

```
cutree(clustering, k = 2)
```

[illegible]

## Principle Component Analysis (PCA)

## PCA of U.K food

First we will read the provided `UK_foods.csv` input file (note we can read this directly from the following tinyurl short link: “<https://tinyurl.com/UK-foods>”).

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66

Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
#dim(x)
dim(x)
```

```
[1] 17  4
```

```
#we could also use nrow(x); ncol(x)
```


We can use `dim(x)` to see how many rows and columns are in our data set.

**Q2.** Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

making the function `row.names` equal to 1 is more efficient and allows us and RStudio to read the data properly.

Now we can generate some basic visuals

```
barplot(as.matrix(x), col= rainbow(nrow(x)), beside=T)
```




**Q3:** Changing what optional argument in the above **barplot()** function results in the following plot?

Setting **beside=FALSE** would make it so the data is not side by side but instead vertically plotted like below.


```
barplot(as.matrix(x), col= rainbow(nrow(x)), beside=F)
```





Let's refine our barplot...

```
pairs(x, col= rainbow(nrow(x)), pch= 16)
```



**Q5:** Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

The diagonal line represents the distribution of the variable, that the variable plotted on x is the same as the variable plotted on y.

**Q6.** What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

N. Ireland seems to consume the most food overall compared to the three other listed nations.

## Applying PCA

Let's apply PCA using the command `prcomp()`. This function expects the transpose of our data.

```
pca <- prcomp (t(x))
summary(pca)
```


Importance of components:

PC1      PC2      PC3      PC4

Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's plot the PCA results

```
plot(pca)
```




We need to access the results of the PCA exploring the dataframe.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

Plotting:

```
plot( x=pca$x[,1], y=pca$x[,2] )
```




**Q7.** Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

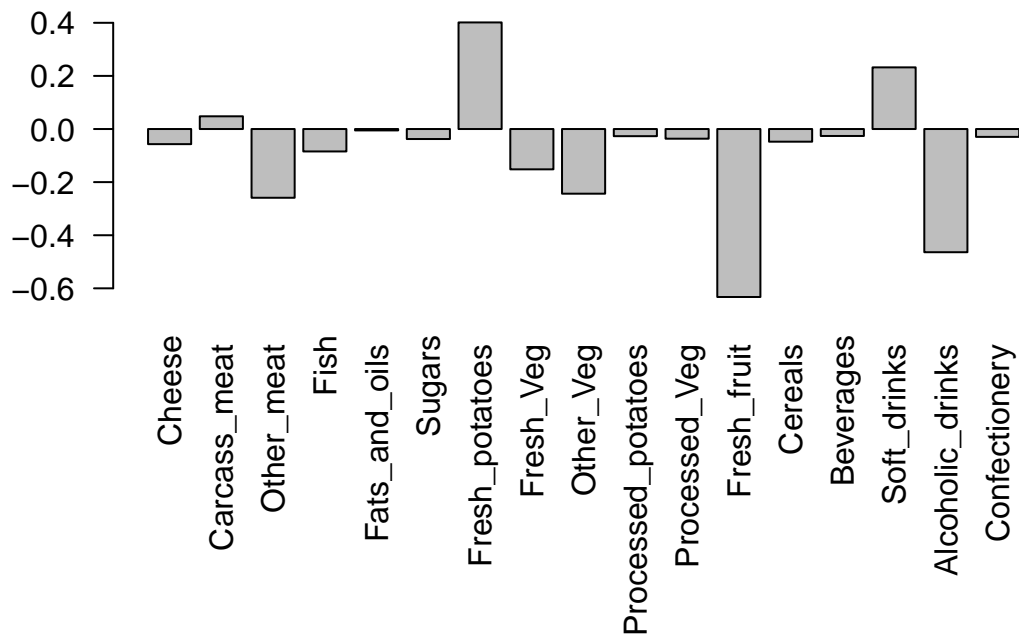
**Q8.** Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

Plotted below:

```
plot( x=pca$x[,1], y=pca$x[,2])  
countrycolor <- c("orange", "pink","blue","green" )  
text( x=pca$x[,1], y=pca$x[,2], colnames(x), col= countrycolor )
```




```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



**Q9:** Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot(pca$rotation[,2], las=2)
```



The most prominent groups in PCA 2 were fresh potatoes and soft drink, which PCA 2 tells us about the second largest amount of variability second to PCA 1.

## PCA of RNA- Sequence dataset

We first load in the file:

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

**Q10:** How many genes and samples are in this data set?

```
dim(rna.data)
```

```
[1] 100  10
```


There are 100 genes, and 10 samples.

```
## Again we have to take the transpose of our data
```

```
pca_rna <- prcomp(t(rna.data), scale=TRUE)
```

```
## Simple un polished plot of pc1 and pc2
```

```
plot(pca_rna$x[,1], pca_rna$x[,2], xlab="PC1", ylab="PC2")
```




```
## Again we have to take the transpose of our data
```

```
pca <- prcomp(t(rna.data), scale=TRUE)
```


```
## Simple un polished plot of pc1 and pc2
```

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```





```
plot(pca, main="Quick scree plot")
```



```
## Variance captured per PC
pca.var <- pca$sdev^2


## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
      xlab=paste0("PC1 (", pca.var.per[1], "%)"),
      ylab=paste0("PC2 (", pca.var.per[2], "%)"))


text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```



```
library(ggplot2)

df <- as.data.frame(pca$x)

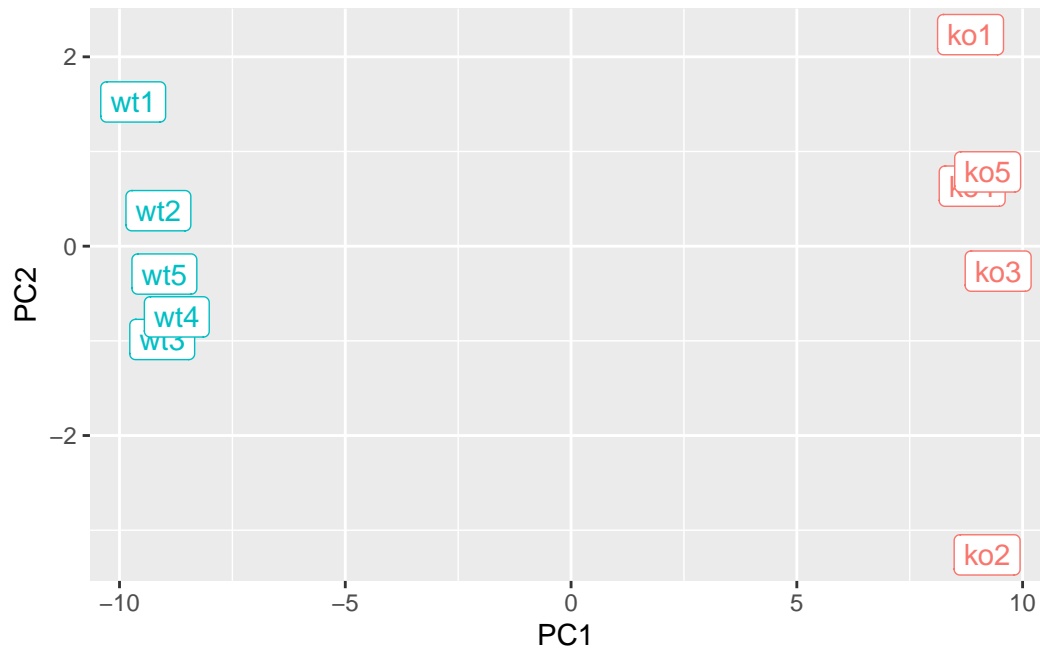
# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)


p
```



```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clearly separates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="Class example data") +
  theme_bw()
```

## PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



Class example data